Group Debate on AI Ethics, Bias, Fairness, Responsible AI

DEBATE TRANSCRIPT

OPENING STATEMENTS

Speaker 1 (FOR – Moderator/Opener,):
Good morning. High-stakes AI decisions — medical diagnosis, criminal sentencing, loan approvals, autonomous vehicles — directly affect human lives, liberty, and livelihood. Our position is clear: human oversight must be mandatory in such decisions, because AI lacks moral accountability, operates as a 'black box,' and has repeatedly failed in the real world. The EU AI Act (2024) itself mandates human oversight for high-risk systems. We argue this is not a brake on technology — it is a safeguard for humanity.

Speaker 5 (AGAINST – Opener,):
We respectfully disagree. Mandating human oversight assumes humans are better decision-makers — but evidence says otherwise. Studies show human reviewers suffer from automation bias — blindly accepting AI outputs (Skitka et al., 2000). Radiologists using AI assistance miss 46% more findings when unsupervised re-checking is removed. Mandatory oversight adds cost, delay, and a false sense of accountability. The right approach is rigorous auditing and validation of AI systems, not ceremonial human sign-off.

MAIN ARGUMENTS — FOR THE MOTION)

Speaker 2 (FOR – Ethical Framework: Accountability):
Under deontological ethics, humans are moral agents; an algorithm cannot be held responsible. In the Dutch childcare benefits scandal (2019–2021), an AI system falsely labelled over 20,000 families as fraudsters, ruining lives — the government resigned because *no human checked* the outputs before enforcement. If a human caseworker had mandated review rights, the harm would have been caught. Accountability requires a human in the loop.

Speaker 3 (FOR – Evidence: Algorithmic Bias):
AI inherits and amplifies bias. The COMPAS recidivism tool (ProPublica, 2016) flagged Black defendants as high-risk at nearly twice the rate of white defendants who did not reoffend. In

healthcare, Obermeyer et al. (Science, 2019) found an algorithm used on 200 million US patients underestimated the illness of Black patients because it used cost as a proxy for need. Judges and doctors — with human oversight — can catch and correct these failures before harm occurs.

Speaker 4 (FOR – Transparency & Black Box Problem):
Generative AI models like GPT-class systems are opaque — even their creators cannot fully explain outputs. The 'right to explanation' in the GDPR (Article 22) grants individuals the right not to be subject to purely automated decisions with significant effects. Human oversight operationalises this right. In high-stakes domains like nuclear plants and aviation, human override has saved countless lives — AI should be held to the same standard.

MAIN ARGUMENTS — AGAINST THE MOTION

Speaker 6 (AGAINST – Human Fallibility & Automation Bias):
Humans are neither neutral nor reliable overseers. Research on automation bias shows humans defer to machine output even when it is wrong — in one study, operators approved 84% of algorithmically incorrect recommendations. A rubber-stamping human is worse than no human: it manufactures the *appearance* of accountability. Better to invest in model validation, bias audits, and explainability (e.g., SHAP/LIME) than mandatory sign-off.

Speaker 7 (AGAINST – Evidence AI Outperforms Humans in Domains):
AI consistently matches or exceeds experts in narrow high-stakes tasks. Esteva et al. (Nature, 2017) showed dermatology-level skin cancer classification from images. IDx-DR (FDA-approved, 2018) autonomously detects diabetic retinopathy with 87% sensitivity — expanding screening to clinics with *no* ophthalmologist, where mandatory human oversight would simply mean no screening at all. Requiring a doctor's signature would defeat the very purpose of equitable access.

Speaker 8 (AGAINST – Practicality, Speed & Cost):
In fraud detection, 100 milliseconds matter — blocking a card before the fraudulent transaction completes. Human review of millions of transactions per day is physically impossible; forcing it would cripple security. Visa/Mastercard-style hybrid systems work precisely *because* AI acts first and humans handle only flagged edge cases. The correct principle is risk-proportionate oversight — not blanket mandates that add latency, cost, and reduce accuracy (studies show review fatigue degrades human accuracy by up to 30%).

REBUTTALS

Rebuttal — FOR (Speaker 3):

Our opponents say AI outperforms doctors — but Esteva's results were later challenged for dataset bias, and IDx-DR still fails on certain retinal conditions. 'AI works when it works' is not a safety policy. The EU AI Act, NIST AI RMF (2023), and WHO guidance on AI in health all converge on one principle: high-risk AI requires human oversight. Are all these global bodies wrong?

Rebuttal — AGAINST (Speaker 7):

We never said *zero* oversight — we oppose mandatory blanket oversight. The EU Act itself allows oversight to be proportionate and can be bypassed where the system is validated as safe. And on automation bias: oversight doesn't prevent rubber-stamping; audits and independent evaluation do. Speaker 3's own example — COMPAS — was used *with* human judges. Oversight existed; bias harm happened anyway. Oversight is necessary but demonstrably insufficient.

CONCLUDING STATEMENTS

Speaker 4 (FOR – Conclusion):

The strongest argument for the motion is accountability: only a human can be a moral agent, and the Dutch benefits scandal proves what happens when automation runs unchecked. Where lives, liberty, and livelihoods are at stake, the burden of proof must lie with the machine.

Speaker 8 (AGAINST – Conclusion):

The strongest argument against is effectiveness: mandatory oversight creates automation bias, cost, and delay while blocking life-saving autonomous AI like IDx-DR. The future is risk-tiered, audit-based governance — human accountability for design and deployment, not human rubber-stamping of every decision.

Speaker 1 (Moderator – Synthesis):

Both sides agree AI must be governed; they differ on *where* the human sits — in the loop (FOR) or on top of the loop through audits (AGAINST). The synthesis emerging from the EU AI Act is 'meaningful human oversight, proportionate to risk' — the strongest ground from both perspectives.